

Tidyverse Create Assignment

Noori Selina

2023-11-07

```
# Load necessary libraries  
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --  
## v dplyr      1.1.3      v readr      2.1.4  
## v forcats    1.0.0      v stringr   1.5.0  
## v ggplot2    3.4.4      v tibble    3.2.1  
## v lubridate  1.9.3      v tidyr     1.3.0  
## v purrr      1.0.1  
## -- Conflicts ----- tidyverse_conflicts() --  
## x dplyr::filter() masks stats::filter()  
## x dplyr::lag()     masks stats::lag()  
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

Introduction

This vignette demonstrates how to perform data manipulation and analysis using the dplyr package. We will work with the Census Income dataset, which contains information on individuals' demographics and income levels. The data set can be found here: <https://www.kaggle.com/datasets/tawfikelmetwally/census-income-dataset/data>

Data Loading

To load the data, I will be importing the data from a GitHub link.

```
census_data <- read_csv("https://raw.githubusercontent.com/NooriSelina/Data-607/main/censusincome.csv")
```

```
## Rows: 32561 Columns: 15  
## -- Column specification -----  
## Delimiter: ","  
## chr (9): Workclass, Education, Marital Status, Occupation, Relationship, Rac...  
## dbl (6): Age, Final Weight, EducationNum, Capital Gain, capital loss, Hours ...  
##  
## i Use 'spec()' to retrieve the full column specification for this data.  
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
head(census_data)
```

```
## # A tibble: 6 x 15
##   Age Workclass      'Final Weight' Education EducationNum 'Marital Status'
##   <dbl> <chr>          <dbl> <chr>          <dbl> <chr>
## 1   39 State-gov          77516 Bachelors          13 Never-married
## 2   50 Self-emp-not-inc    83311 Bachelors          13 Married-civ-spou~
## 3   38 Private           215646 HS-grad            9 Divorced
## 4   53 Private           234721 11th             7 Married-civ-spou~
## 5   28 Private           338409 Bachelors          13 Married-civ-spou~
## 6   37 Private           284582 Masters          14 Married-civ-spou~
## # i 9 more variables: Occupation <chr>, Relationship <chr>, Race <chr>,
## #   Gender <chr>, 'Capital Gain' <dbl>, 'capital loss' <dbl>,
## #   'Hours per Week' <dbl>, 'Native Country' <chr>, Income <chr>
```

Data Manipulation with dplyr

The dplyr package provides a set of functions for data manipulation. We will use the following functions to explore the dataset.

1. Filtering - Using the dplyr package, I will use the `filter()` function to filter the dataset based on specific criteria. In this case, we are interested in individuals with incomes greater than \$50,000, so we filter the dataset to include only such individuals.

```
high_income_data <- census_data %>%
  filter(Income == ">50K")
head(high_income_data)
```

```
## # A tibble: 6 x 15
##   Age Workclass      'Final Weight' Education EducationNum 'Marital Status'
##   <dbl> <chr>          <dbl> <chr>          <dbl> <chr>
## 1   52 Self-emp-not-inc    209642 HS-grad            9 Married-civ-spo~
## 2   31 Private           45781 Masters          14 Never-married
## 3   42 Private           159449 Bachelors          13 Married-civ-spo~
## 4   37 Private           280464 Some-coll~    10 Married-civ-spo~
## 5   30 State-gov          141297 Bachelors          13 Married-civ-spo~
## 6   40 Private           121772 Assoc-voc       11 Married-civ-spo~
## # i 9 more variables: Occupation <chr>, Relationship <chr>, Race <chr>,
## #   Gender <chr>, 'Capital Gain' <dbl>, 'capital loss' <dbl>,
## #   'Hours per Week' <dbl>, 'Native Country' <chr>, Income <chr>
```

2. Grouping and Summarizing - The `group_by()` and `summarize()` functions of the dplyr package are valuable for aggregating data. We will group the data by education level and calculate summary statistics for age and hours worked per week within each education category.

```
income_summary <- high_income_data %>%
  group_by(Income) %>%
  summarize(
    mean_age = mean(Age, na.rm = TRUE),
    median_hours = median('Hours per Week', na.rm = TRUE)
```

```
)

print(income_summary)

## # A tibble: 1 x 3
##   Income mean_age median_hours
##   <chr>      <dbl>      <dbl>
## 1 >50K      44.2        40
```

3. Sorting - The `arrange()` function is used to sort the summarized data. In this example, we sort the education summary by mean age in descending order, which helps identify the education categories with the highest mean age.

```
high_income_data <- high_income_data %>%
  arrange(desc(Age))

print(head(high_income_data, 10))
```

```
## # A tibble: 10 x 15
##   Age Workclass   'Final Weight' Education EducationNum 'Marital Status'
##   <dbl> <chr>      <dbl> <chr>      <dbl> <chr>
## 1  90 Local-gov  227796 Masters      14 Married-civ-spou~
## 2  90 Private   51744 Masters      14 Never-married
## 3  90 Private   87372 Prof-school 15 Married-civ-spou~
## 4  90 Private   46786 Bachelors    13 Married-civ-spou~
## 5  90 Private  175491 HS-grad     9 Married-civ-spou~
## 6  90 Private   88991 Masters    13 Married-civ-spou~
## 7  90 Private  206667 Masters    14 Married-civ-spou~
## 8  90 ?        313986 HS-grad     9 Married-civ-spou~
## 9  84 Self-emp-inc 172907 Some-college 10 Married-civ-spou~
## 10 83 Self-emp-inc 240150 10th        6 Married-civ-spou~
## # i 9 more variables: Occupation <chr>, Relationship <chr>, Race <chr>,
## #   Gender <chr>, 'Capital Gain' <dbl>, 'capital loss' <dbl>,
## #   'Hours per Week' <dbl>, 'Native Country' <chr>, Income <chr>
```

Conclusion

In this vignette, the `dplyr` package from the `tidyverse` collection was used to efficiently manage and analyze the Census Income dataset. Specifically, the data was filtered using the `filter()` function to focus on high-income individuals, grouped by income levels using the `group_by()` function, and then sorted by mean age with the `arrange()` function.

This approach uncovered valuable insights about the demographics of high-income individuals. The `dplyr` package, along with other tools from the `tidyverse`, made data manipulation tasks straightforward and effective.